

Introduction to Neural Computation

Prof. Michale Fee
MIT BCS 9.40 — 2018

Lecture 14
Rate models and Perceptrons

Game plan for Lectures 14 – 18

Examine the computational properties of networks of neurons

- Rate models
- Feed-forward neural networks (Perceptrons)
- Matrix operations
- Basis sets
- Principal components analysis
- Recurrent Neural Networks
- Line attractors in short term memory
- Hopfield networks

Learning Objectives for Lecture 14

- Derive a mathematically tractable model of neural networks
(the rate model)
- Building receptive fields with neural networks
- Vector notation and vector algebra
- Neural networks for classification
- Perceptrons

Neural Network Models

- We are going to examine some of the computational properties of networks of neurons.
- Our first step is to derive a simplified mathematical model of neurons that we can study analytically.

Why would we want to do this?

- For example, we approximated the detailed spiking properties of neurons (HH model) with an integrate & fire (IF) model.
- This is not enough of a simplification to develop an analytical model of neural circuits.

Rate models

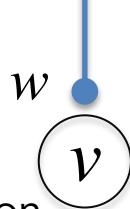
- Let's start with two neurons, an input neuron that synapses with weight w onto an output neuron.
- We are going to ignore spike times, and describe the inputs and outputs of our neurons simply as firing rates.
- In the simplest case... linear neurons! $v = wu$

How can we justify this?

input neuron



u = firing rate of input neuron



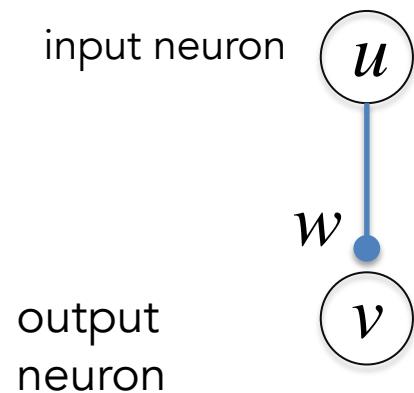
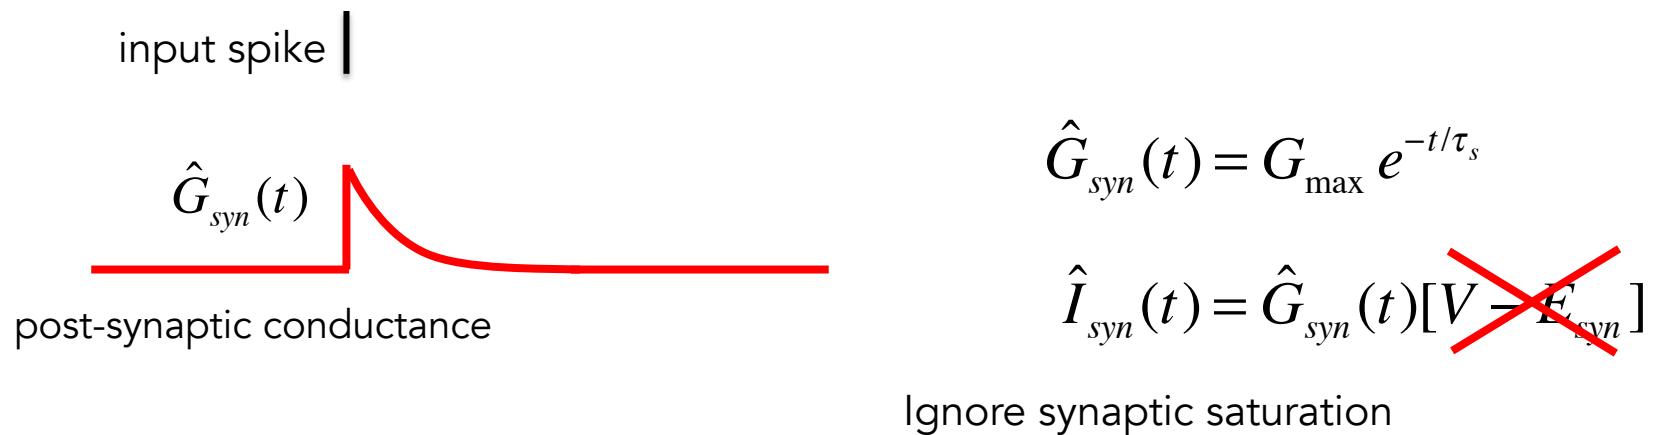
w = synaptic strength (weight)

v = firing rate of output neuron

output neuron

Rate models

- Let's examine the response of the output neuron to a single input spike.



Now write $\hat{I}_{syn}(t) = wK(t)$

where $K(t)$ is a kernel of unit area

$$K(t) = \frac{1}{\tau_s} e^{-t/\tau_s} \quad \text{area} = 1$$

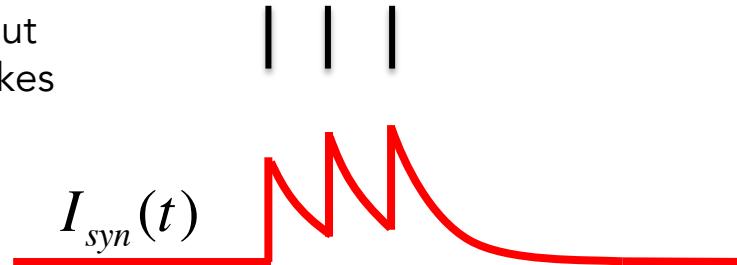
Rate models

- How do we get the response to multiple input spikes?

Input spike train

$$\rho(t) = \sum_i \delta(t - t_i)$$

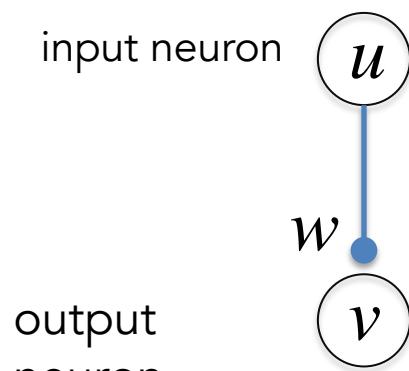
input
spikes



- We convolve the input spike train with the synaptic kernel !

$$I_{syn}(t) = w K * \rho(t)$$

But what is this?



- If K is a kernel with area normalized to one, then...

$$K * \rho(t) = u(t)$$

is just the firing rate of the input neuron!

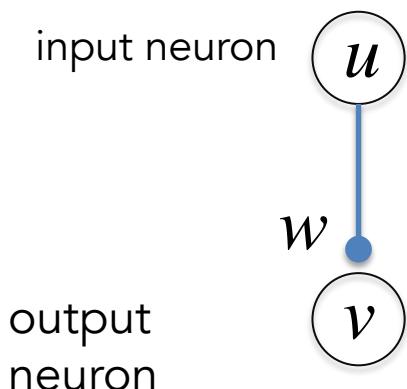
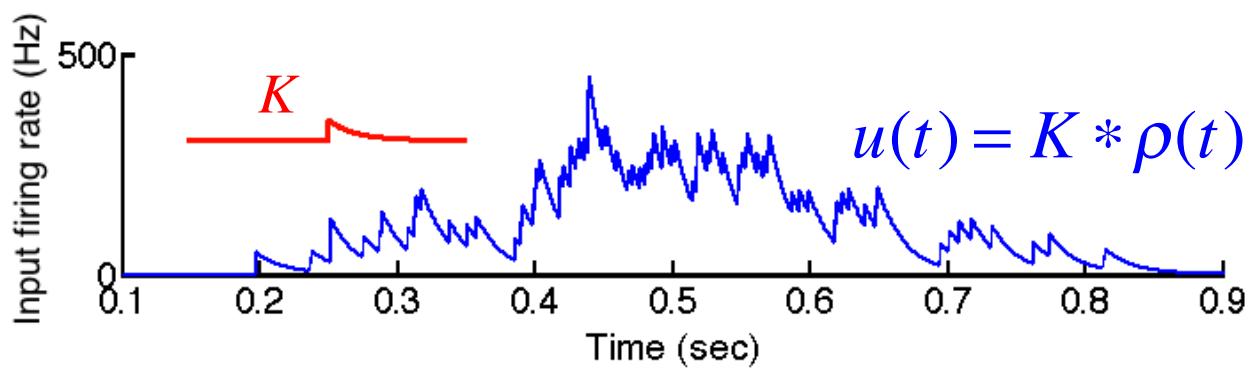
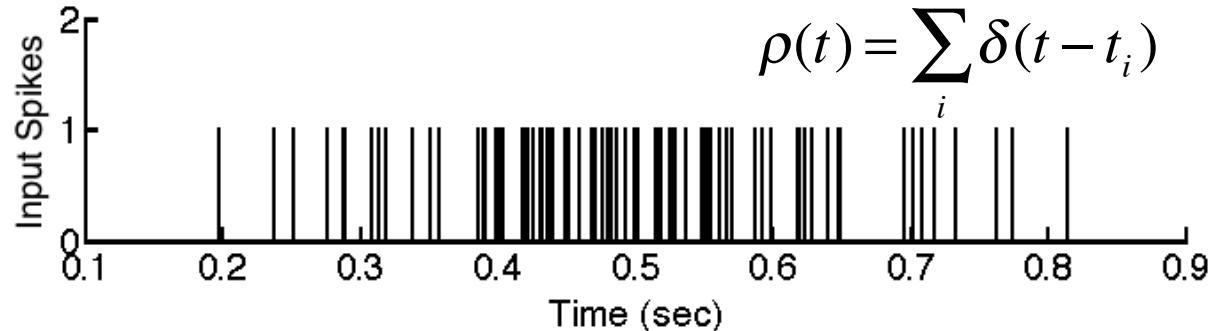
Rate models

- Thus, we can write presynaptic firing rate as

`pulse_of_firing_spikes.m`

$$u(t) = K * \rho(t)$$

$$\rho(t) = \sum_i \delta(t - t_i)$$

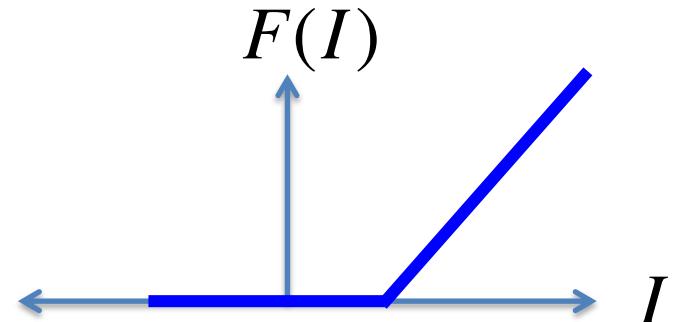
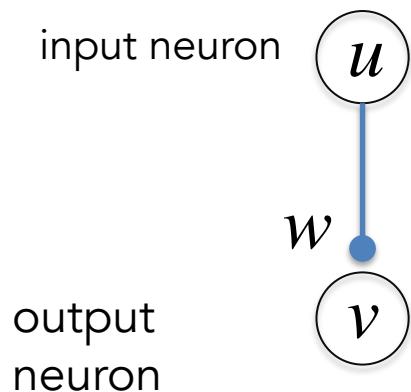


$$\begin{aligned} I_{syn}(t) &= w K * \rho(t) \\ &= w u(t) \end{aligned}$$

Rate models

- Now, how about the firing rate of the output neuron?
- You remember that when you injected a constant current into our Integrate and Fire model...

The steady-state firing rate of the neuron had a threshold current below which the neuron would not spike, and some increasing f.r. above threshold.



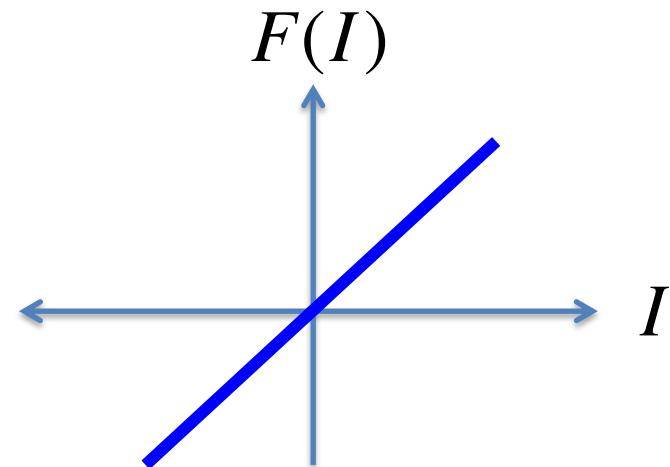
- In general, we can write the output firing rate of our model neuron as:

$$v = F[I_s] = F[wu]$$

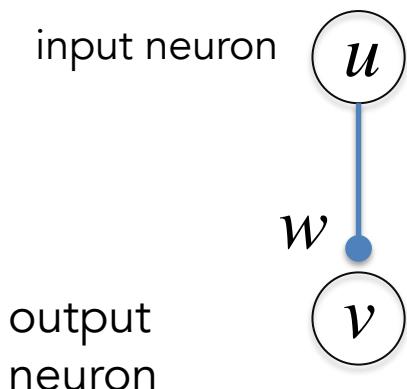
Linear rate models

- We will now consider an even greater simplification of our neurons... assume they are linear.

$$F[x] = x$$



We will come back to non-linear neurons shortly... they will be very important.

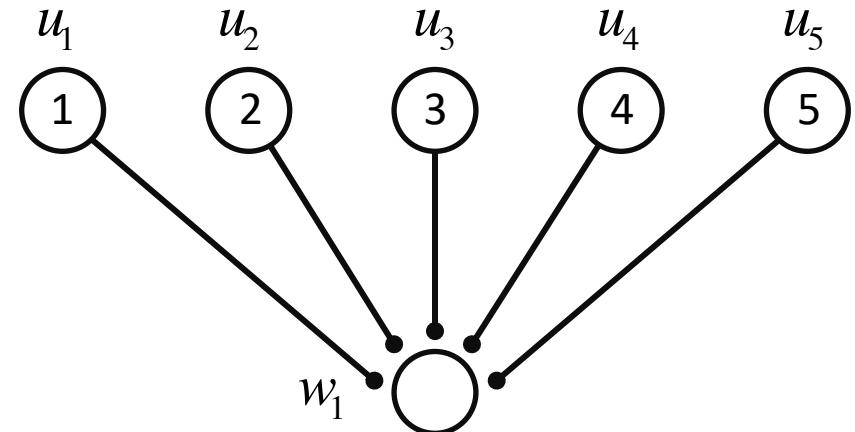


- Of course real neurons can't have negative firing rates, but we can gain a lot of insight using this approximation.
- Thus, we can write the output firing rate of our linear neuron as:

$$v = wu$$

Multiple inputs

- What happens when our output neuron has many inputs?



$$I_{syn} = w_1 u_1 + w_2 u_2 + w_3 u_3 + \dots$$

- The steady-state response of our linear neuron is now:

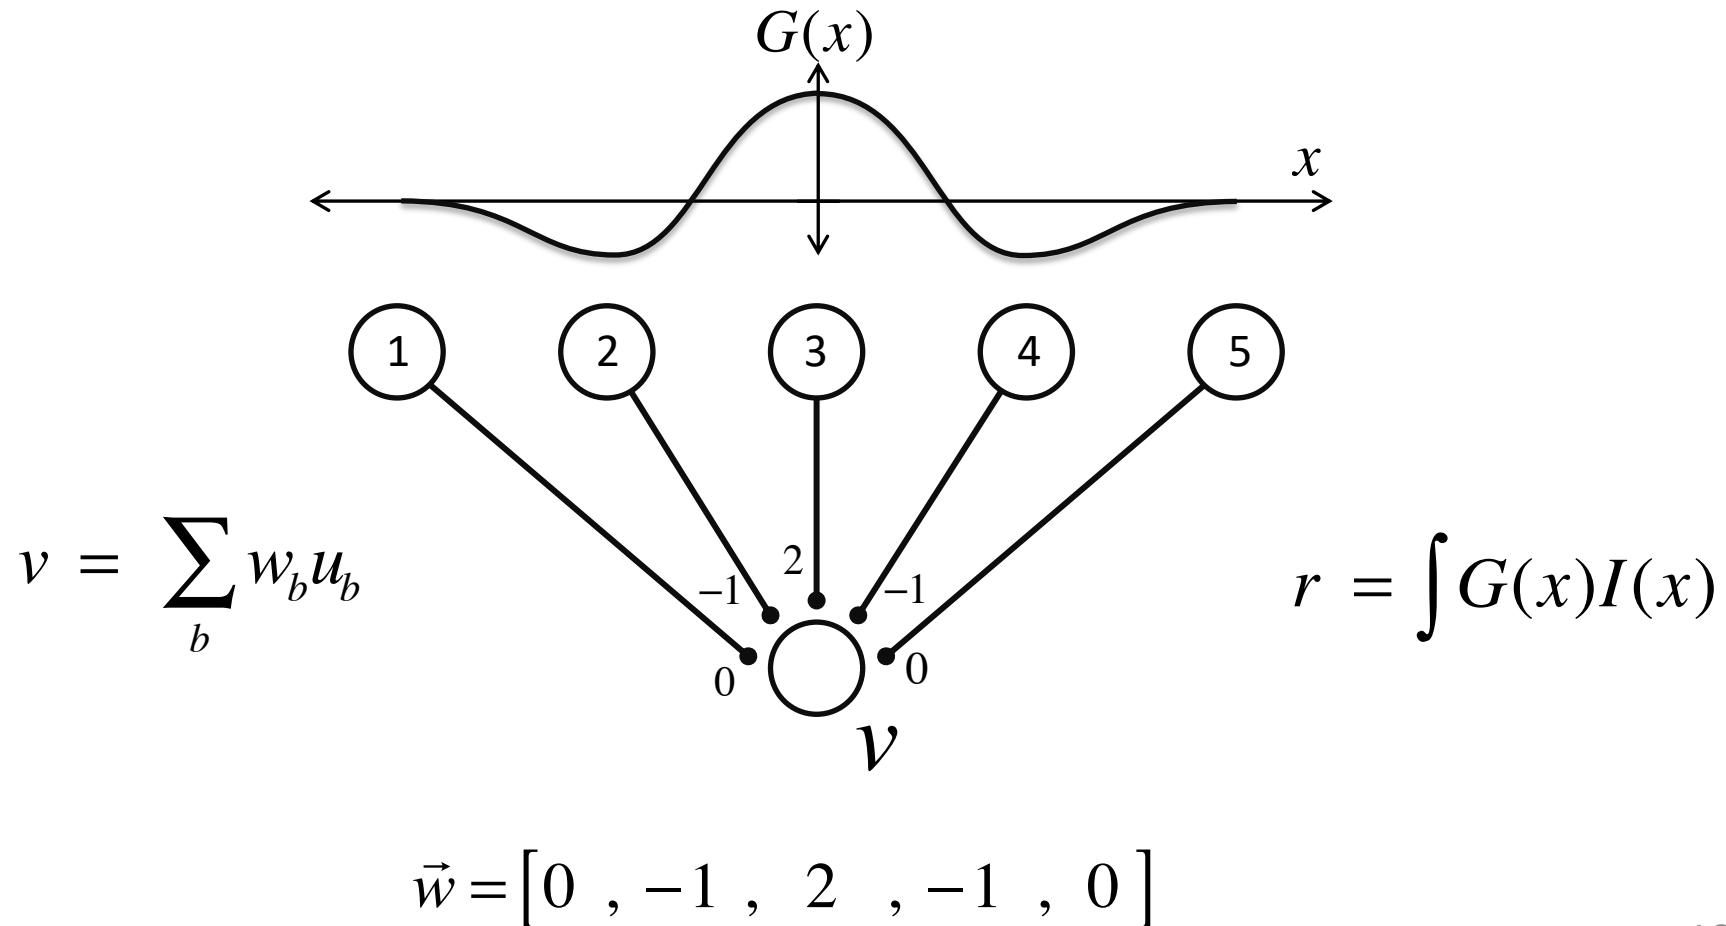
$$v = \sum_b w_b u_b$$

Learning Objectives for Lecture 14

- Derive a mathematically tractable model of neural networks
(the rate model)
- Building receptive fields with neural networks
- Vector notation and vector algebra
- Neural networks for classification
- Perceptrons

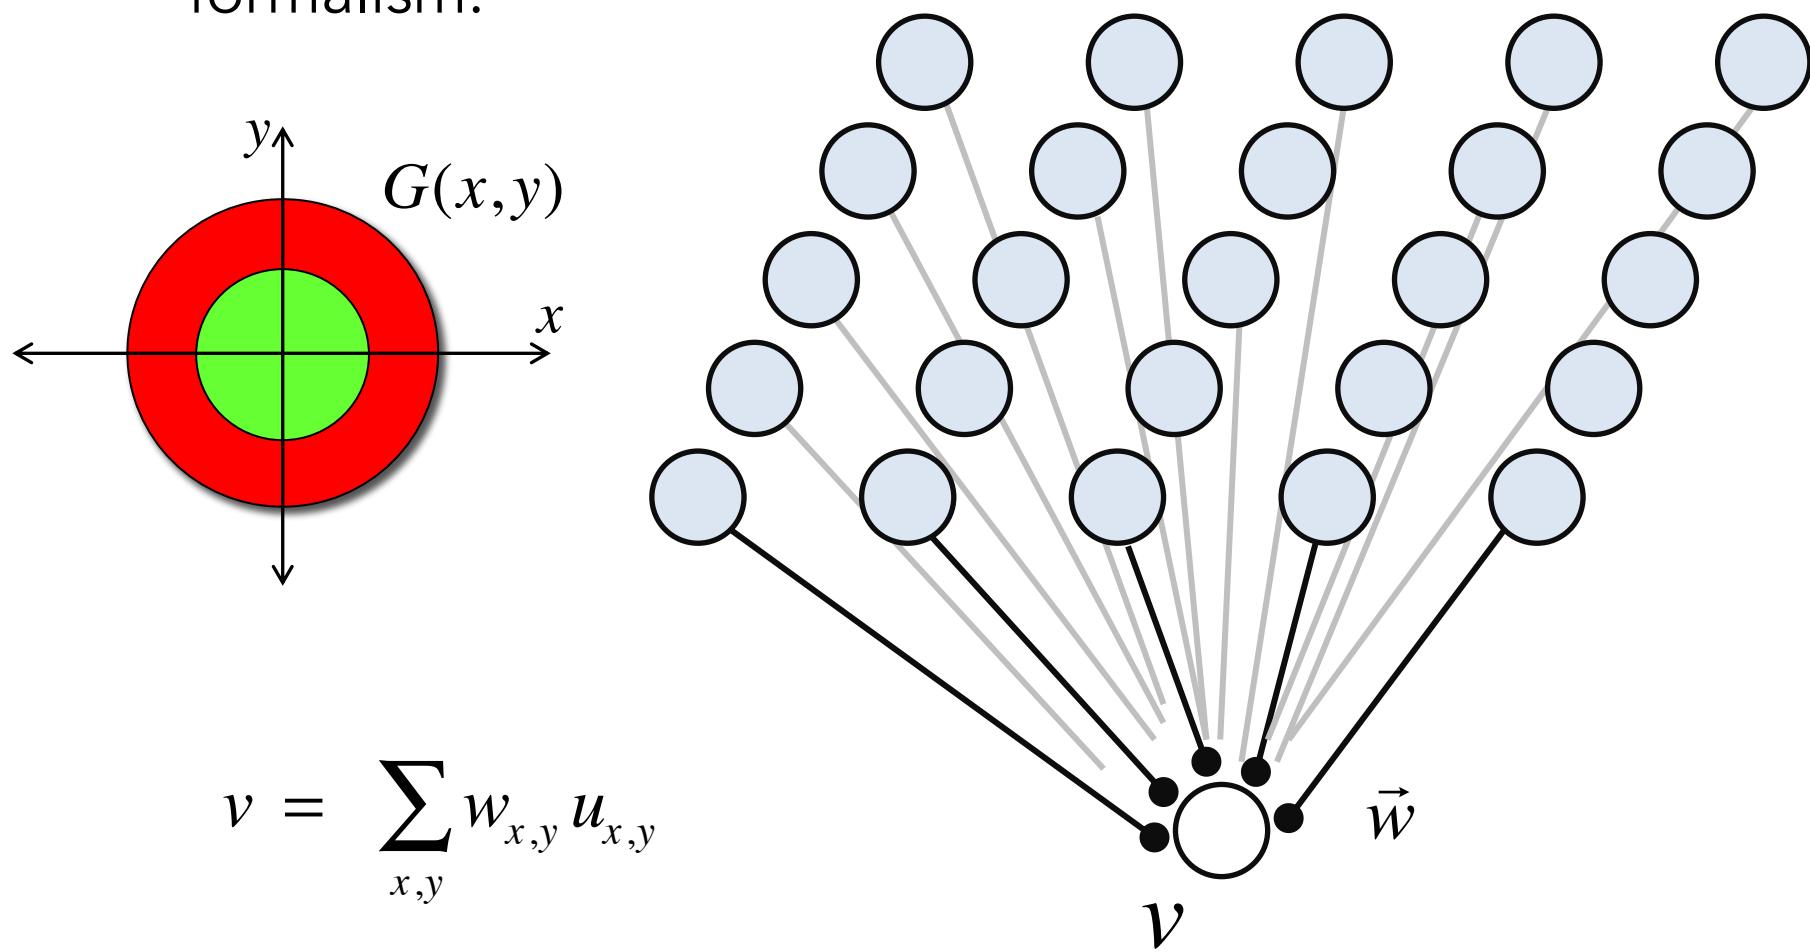
How to build a receptive field

- We can see that the choice of weights allows us to specify the receptive field of our output neuron



How to build a receptive field

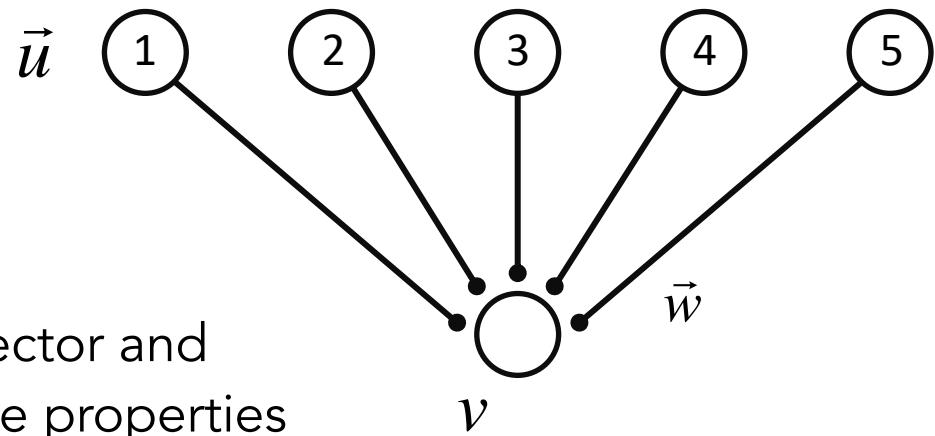
- We can even build 2D receptive fields with the same formalism.



Linear algebra detour

- Mathematically, we have described the response of our linear neuron as

$$v = \sum_b w_b u_b$$



- We are going to start using vector and matrix notation to describe the properties of networks
- ...because it is much more compact and powerful
- We have to take a short detour to learn some linear algebra.

Learning Objectives for Lecture 14

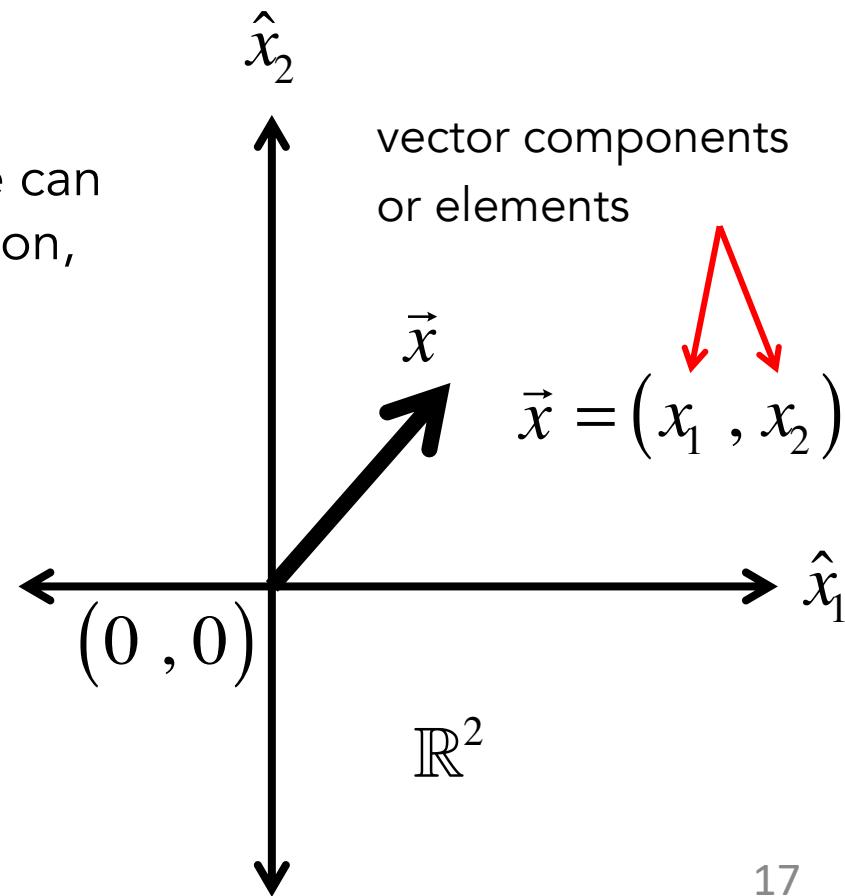
- Derive a mathematically tractable model of neural networks
(the rate model)
- Building receptive fields with neural networks
- Vector notation and vector algebra
- Neural networks for classification
- Perceptrons

Linear algebra detour

- A vector is a collection of numbers.
- The number of numbers in the collection is called the dimensionality of the vector.
- If there are two or three numbers, we can draw a vector as a position, or direction, in space.

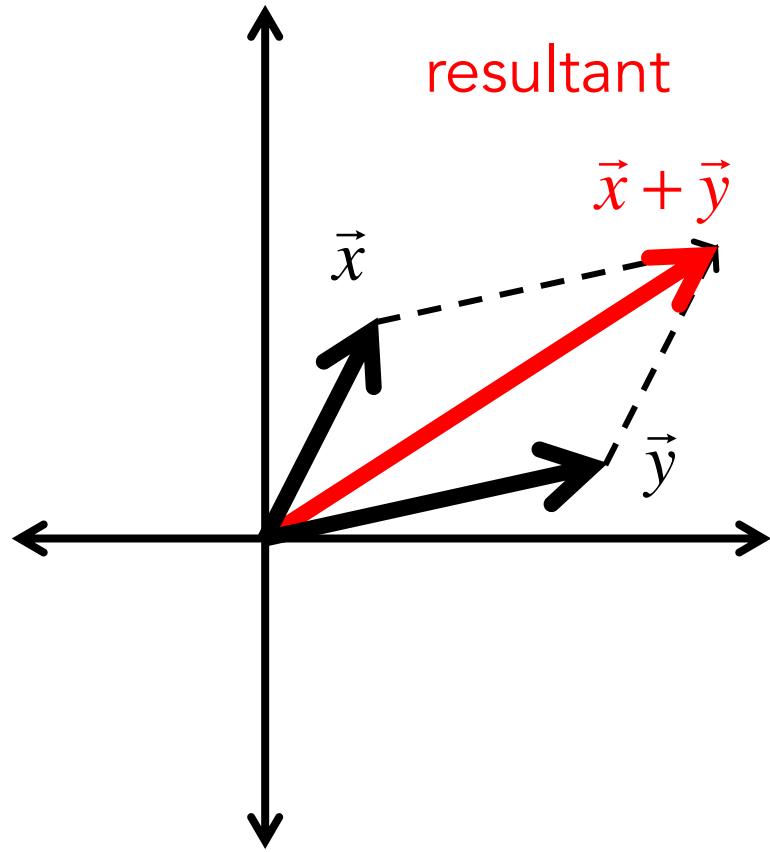
$$\vec{x} = \begin{pmatrix} x_1 & x_2 \end{pmatrix} \quad \text{row vector}$$

$$\vec{x} = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} \quad \text{column vector}$$



Vector sum

- Sum of two vectors



$$\vec{x} + \vec{y} = \begin{pmatrix} x_1 \\ x_2 \\ \dots \\ x_n \end{pmatrix} + \begin{pmatrix} y_1 \\ y_2 \\ \dots \\ y_n \end{pmatrix} = \begin{pmatrix} x_1 + y_1 \\ x_2 + y_2 \\ \dots \\ x_n + y_n \end{pmatrix}$$

Element-by-element
addition

Vector products

- There are several ways of taking the product of two vectors

- Element-by-element product
- Inner product
- Outer product

We will cover this later

- Cross product

Important in physics, but we won't cover this.

Vector products

- Element-by-element product (Hadamard product)

$$\vec{x} \circ \vec{y} = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix} \circ \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix} = \begin{pmatrix} x_1 y_1 \\ x_2 y_2 \\ \vdots \\ x_n y_n \end{pmatrix}$$

- In Matlab, the element-by-element product is `x.*y`

Vector products

- Inner product or dot product

$$\vec{x} = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{pmatrix} \quad \vec{y} = \begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix} \quad \begin{aligned} \vec{x} \cdot \vec{y} &= x_1 y_1 + x_2 y_2 + \dots + x_n y_n \\ &= \sum_{i=1}^n x_i y_i \quad = \text{scalar} \end{aligned}$$

Some properties...

commutative

$$\vec{x} \cdot \vec{y} = \vec{y} \cdot \vec{x}$$

distributive

$$\vec{w} \cdot (\vec{x} + \vec{y}) = \vec{w} \cdot \vec{x} + \vec{w} \cdot \vec{y}$$

linearity

$$(a\vec{x}) \cdot \vec{y} = a(\vec{x} \cdot \vec{y})$$

Vector products

- Inner product in matrix notation

$$\vec{x} \cdot \vec{y} = \begin{matrix} & 1 \times N & & N \times 1 & & 1 \times 1 \\ \vec{x} \cdot \vec{y} = & \left(\begin{array}{cccc} x_1 & x_2 & \cdots & x_n \end{array} \right) & \left(\begin{array}{c} y_1 \\ y_2 \\ \vdots \\ y_n \end{array} \right) & = scalar \end{matrix}$$

- In Matlab...

```
x = [1; 2; 3]; % column vector (1 x 3)  
y = [2; 4; 6]; % column vector (1 x 3)  
z = x' * y; % ' means transpose  
% row*column vector (1x3)*(3x1)
```

Vector products

- Dot product of a vector with itself

$$\vec{x} \cdot \vec{x} = \sum_{i=1}^n x_i x_i = |\vec{x}|^2$$

$|\vec{x}|$ is the 'norm' or 'magnitude' of the vector

$$|\vec{x}| = \sqrt{\sum_{i=1}^n x_i x_i}$$
 Pythagorean theorem

Unit vector

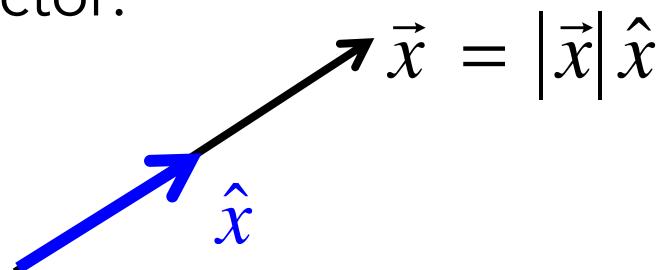
- A unit vector has length 1.

$$|\hat{x}| = 1 \quad \hat{x} \cdot \hat{x} = 1$$

- We can make a unit vector out of any vector

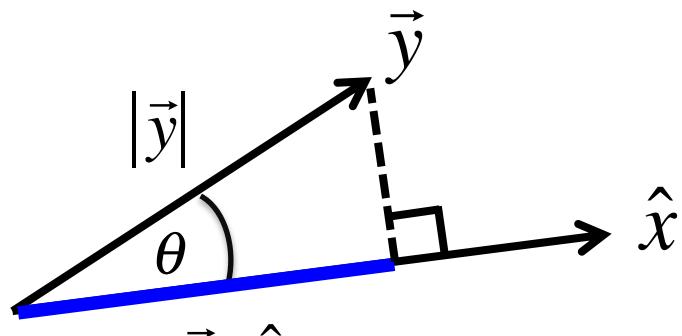
$$\hat{x} = \frac{1}{|\vec{x}|} \vec{x}$$

- We can express any vector as a product of a length times a unit vector:



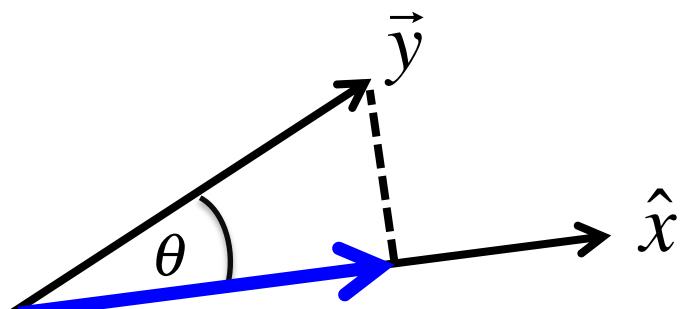
Projection

Find the component of vector \vec{y} in the direction of vector \hat{x} . Let \hat{x} be a unit vector.



$$= |\vec{y}| \cos \theta$$

'Scalar projection' of \vec{y} onto \hat{x}



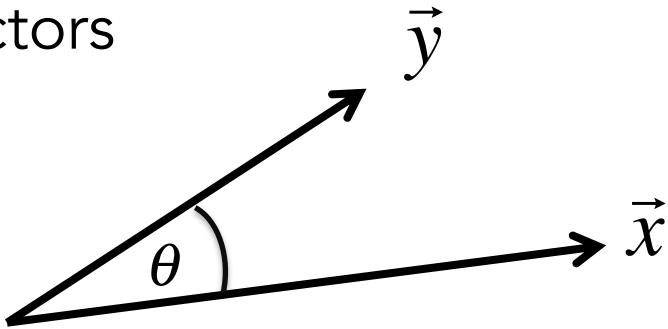
$$(\vec{y} \cdot \hat{x}) \hat{x}$$

= scalar times unit vector

'Vector projection' of \vec{y} onto \hat{x}

Geometric intuition of dot products

- Dot product is related to the cosine of the angle between two vectors



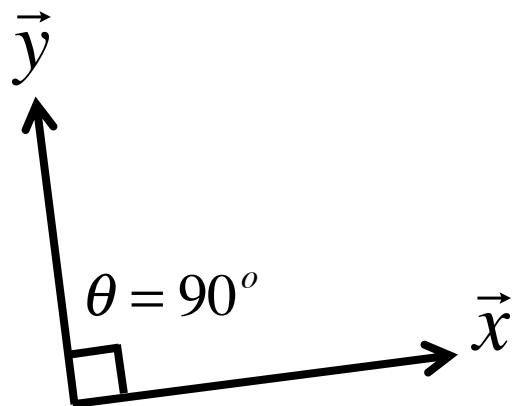
$$\vec{x} \cdot \vec{y} = |\vec{x}| |\vec{y}| \cos \theta \quad \cos \theta = \frac{\vec{x} \cdot \vec{y}}{|\vec{x}| |\vec{y}|}$$

- If x and y are unit vectors, then...

$$\hat{x} \cdot \hat{y} = \cos \theta$$

Orthogonality

- Two vectors are orthogonal (perpendicular) if and only if their dot product is zero.



$$\vec{x} \cdot \vec{y} = |\vec{x}| |\vec{y}| \cos \theta$$

$$\vec{x} \cdot \vec{y} = 0$$

$$\cos 90^\circ = 0$$

- The projection of y onto x is zero.
- The vector projection of y along x is the zero vector.

'Correlation' intuition of dot product

- The dot product is related to the statistical correlation between the elements of the two vectors

$$\cos \theta = \frac{\vec{y} \cdot \vec{x}}{\|\vec{x}\| \|\vec{y}\|} = \frac{\sum_i y_i x_i}{\sqrt{\sum_i x_i x_i} \sqrt{\sum_i y_i y_i}}$$

The diagram shows two blue arrows pointing from the terms in the formula to their respective definitions. One arrow points from the term $\vec{y} \cdot \vec{x}$ to the expression $\sum_i y_i x_i$. Another arrow points from the term $\|\vec{x}\| \|\vec{y}\|$ to the expression $\sqrt{\sum_i x_i x_i} \sqrt{\sum_i y_i y_i}$.

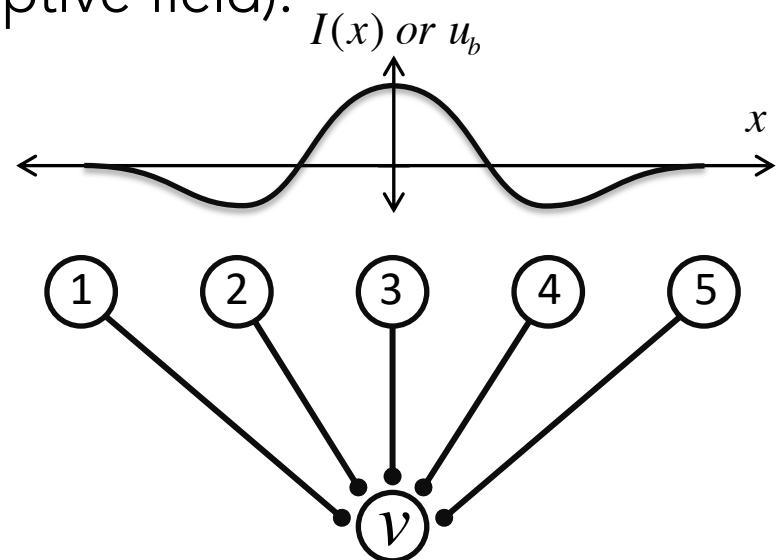
Bounded between -1 and 1

Optimal stimulus

- The response of a neuron is the dot-product of the stimulus vector with the weight vector (receptive field).

$$v = \sum_b w_b u_b$$

$$v = \vec{w} \cdot \vec{u} = |\vec{w}| |\vec{u}| \cos \theta$$



- Thus, for a given amount of power in the stimulus $a^2 = |\vec{u}|^2$, the stimulus that has the best overlap with the receptive field () produces the largest neuronal response.
- We now have a definition of the 'optimal stimulus':

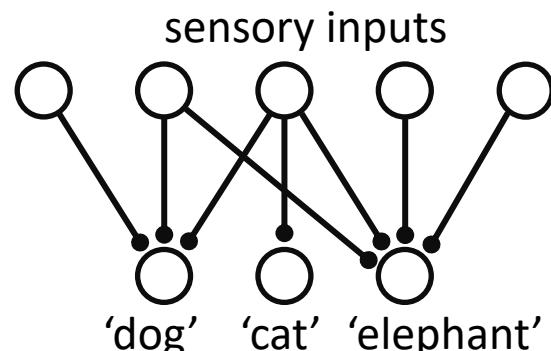
$$\vec{u} = a\hat{w}$$

Learning Objectives for Lecture 14

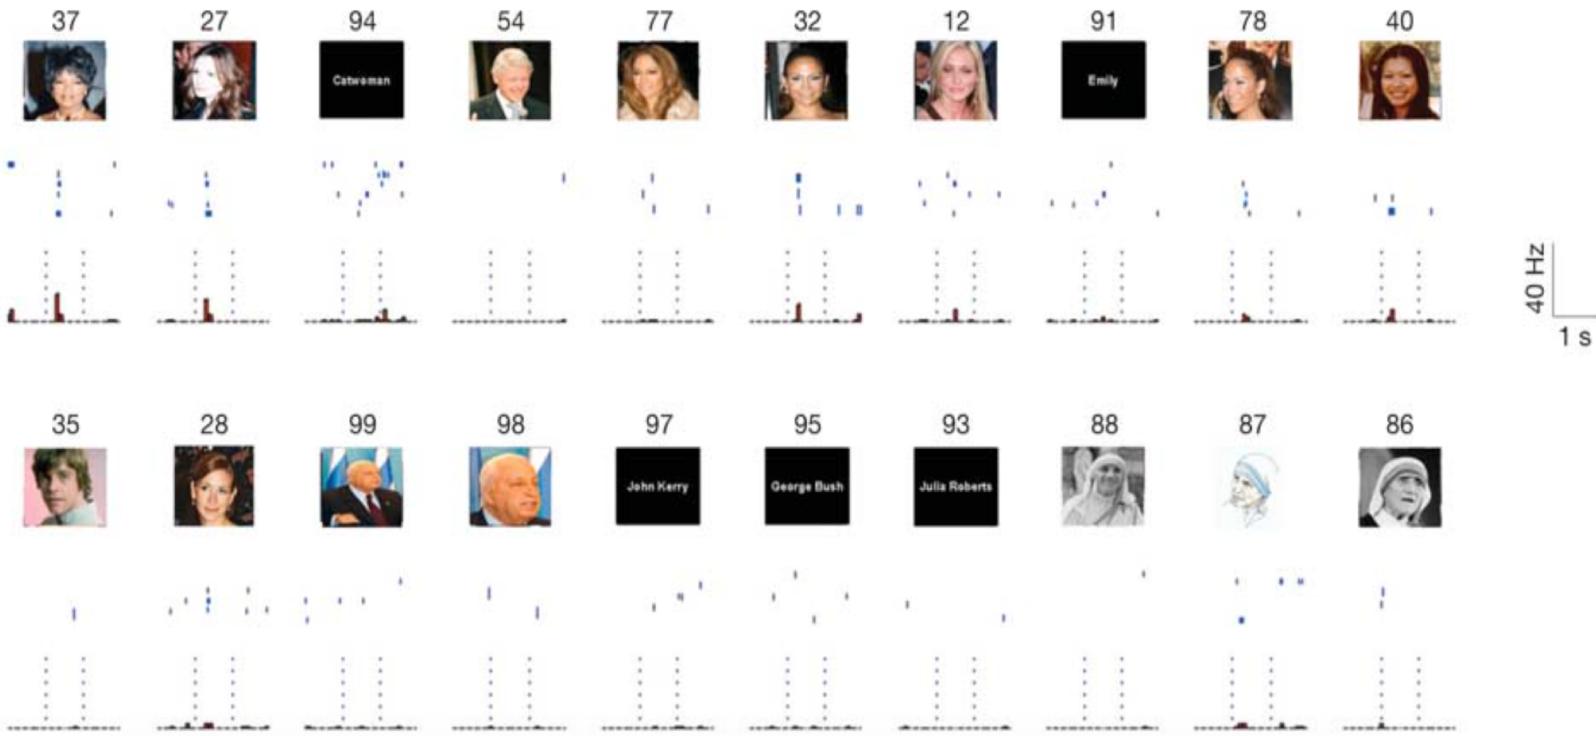
- Derive a mathematically tractable model of neural networks
(the rate model)
- Building receptive fields with neural networks
- Vector notation and vector algebra
- Neural networks for classification
- Perceptrons

Classification

- A general computational problem solved by brain circuits is that of classification.
- Does that visual input represent a house cat or a tiger
 - an edible object or a poisonous one
 - a friendly dog or a wolf
- Feedforward circuits can be very good at classification



Object recognition in human cortex

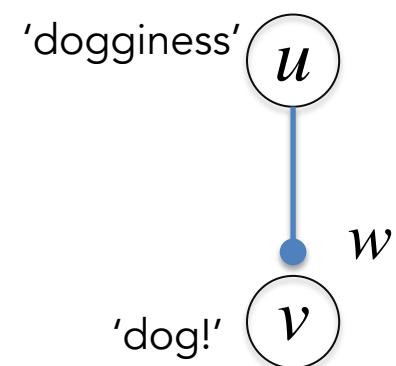
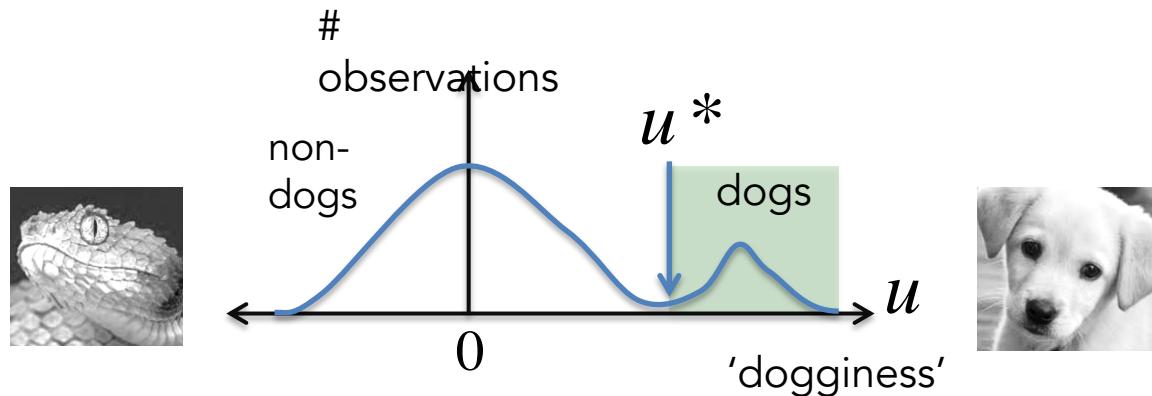


Learning Objectives for Lecture 14

- Derive a mathematically tractable model of neural networks
(the rate model)
- Building receptive fields with neural networks
- Vector notation and vector algebra
- Neural networks for classification
- Perceptrons

Perceptrons

- How do we make a neuron that fires when it sees a dog, but does not fire when there is no dog?
- Classification problem in one dimension: one input neuron whose firing rate is proportional to a **feature** - 'dogginess'.



- A central feature of classification is decision making.
 - There exists a 'classification boundary' in stimulus space that separates dogs from non-dogs.
- How does a neural circuit make a decision? **Spike threshold!**

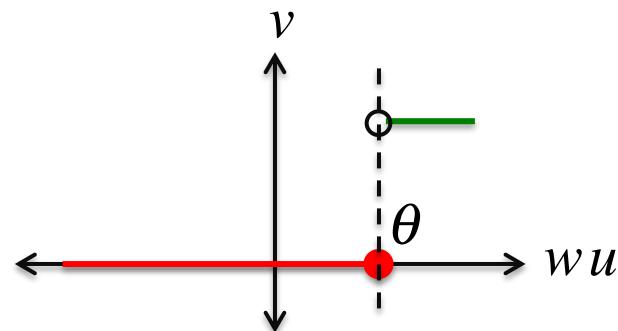
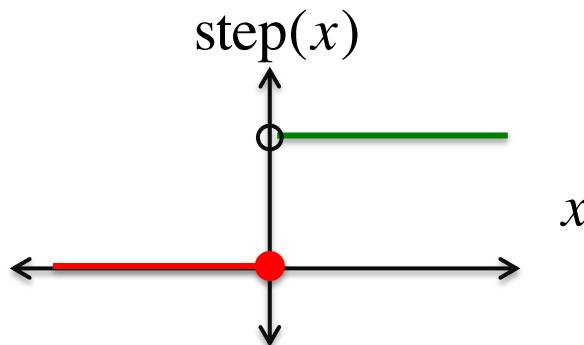
Binary threshold unit

- For a perceptron, we make a simplified model of a neuron that is very good at making decisions:

$$F(x) = \text{step}(x)$$

$$v = F(wu - \theta)$$

Theta is the threshold, not an angle.



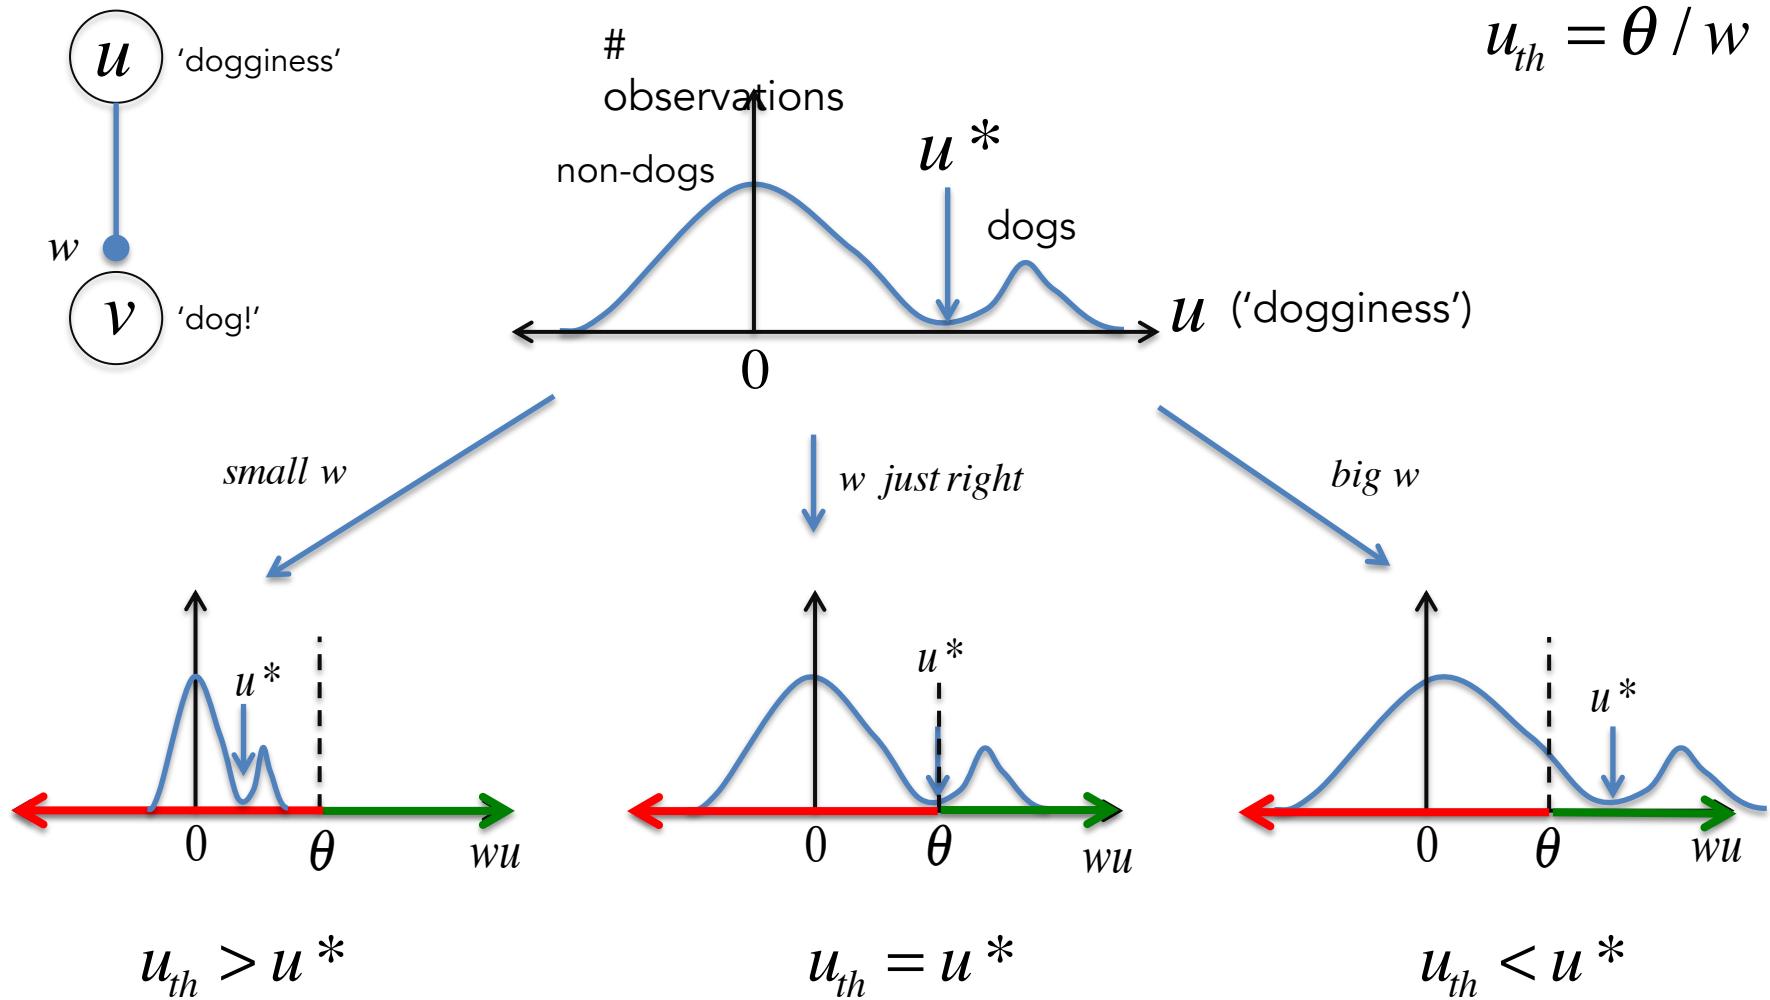
Neuron fires when the input $wu > \theta$

Thus, the output neuron begins to fire when the input neuron has a firing rate greater than the 'decision boundary.'

$$u_{th} = \theta / w$$

Setting the weight

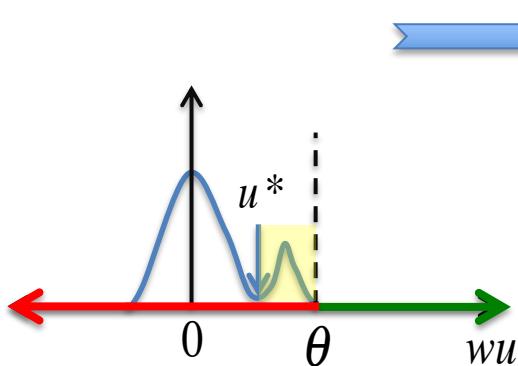
- To classify, we need to learn the right w to make $u_{th} = u^*$



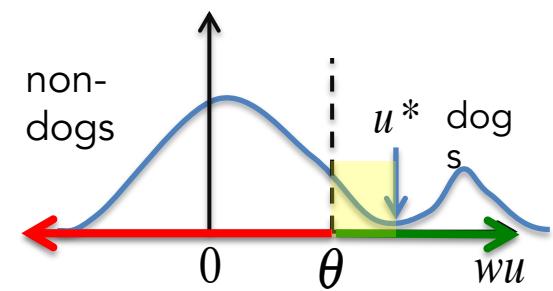
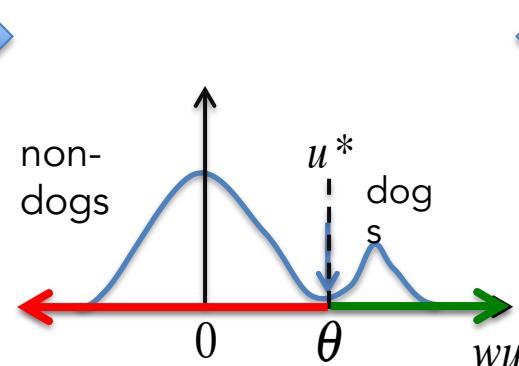
Setting the weight

- To classify, we need to learn the right w to make $u_{th} = u^*$

Error: Dogs classified as non-dogs
⇒ Make w bigger

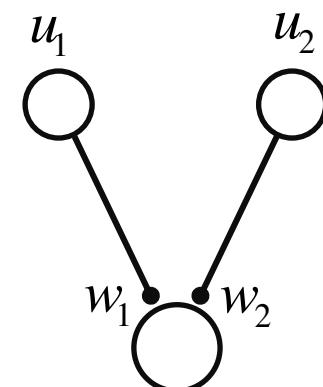
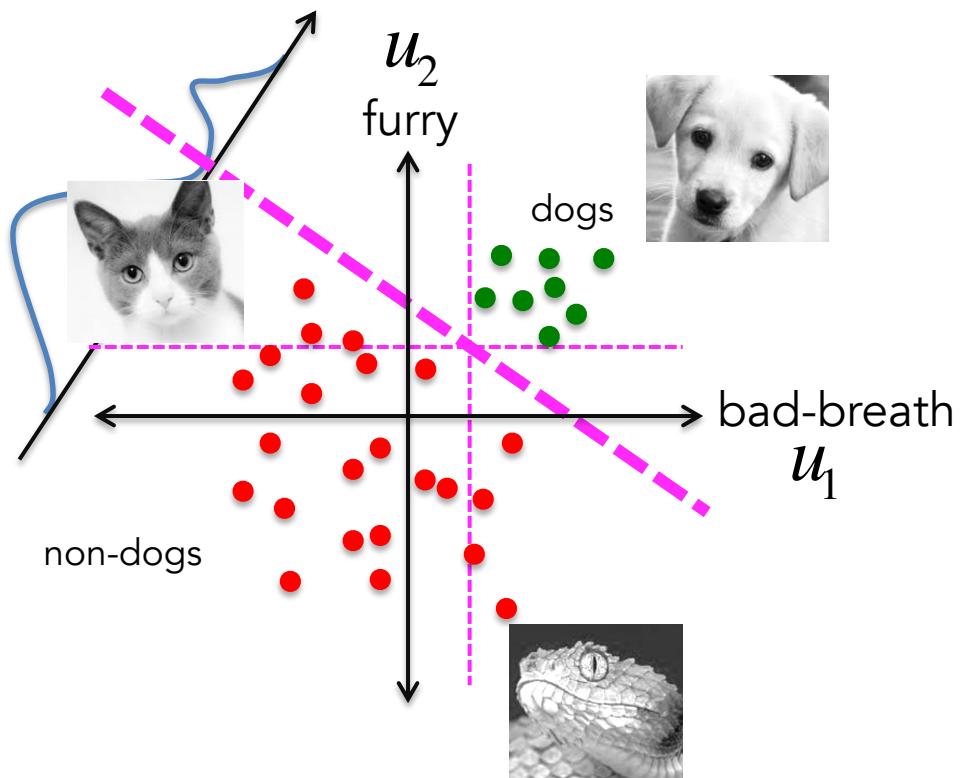


Error: Non-dogs classified as dogs
⇒ Make w smaller



Decision boundary in two dimensions

- Sometimes classification has to be done on the basis of many features, not just one.



$$v = \vec{w} \cdot \vec{u}$$

Decision boundary in two dimensions

- Let's look at the case where our neuron gets two inputs

$$v = F(\vec{w} \cdot \vec{u} - \theta)$$

- Now the decision boundary looks different...

$$\vec{w} \cdot \vec{u} - \theta = 0$$

what is this?

$$\vec{w} \cdot \vec{u} = \theta$$

- This is an equation for a line in the space of \vec{u} , specified by the weights \vec{w} and threshold θ .

$$w_1 u_1 + w_2 u_2 = \theta$$

Decision boundary in two dimensions

- Let's start by looking at the case where $\theta = 0$

$$v = F(\vec{w} \cdot \vec{u})$$

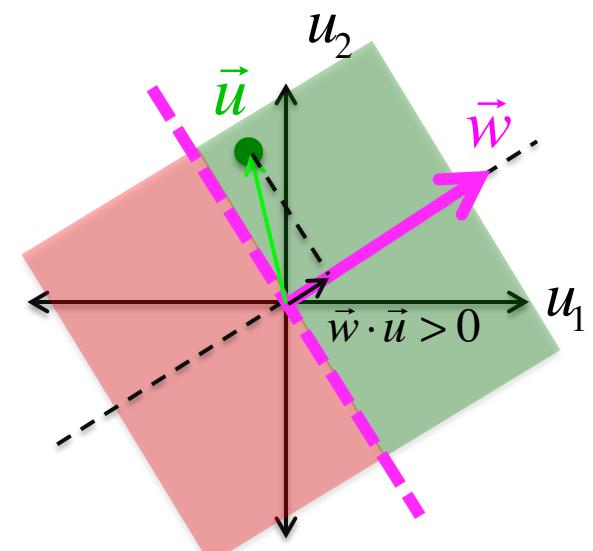
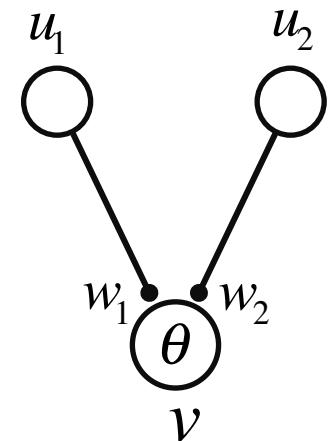
- The neuron now fires when the projection of \vec{u} along \vec{w} is positive $\vec{w} \cdot \vec{u} > 0$

- The decision boundary is given by

$$\vec{w} \cdot \vec{u} = 0$$

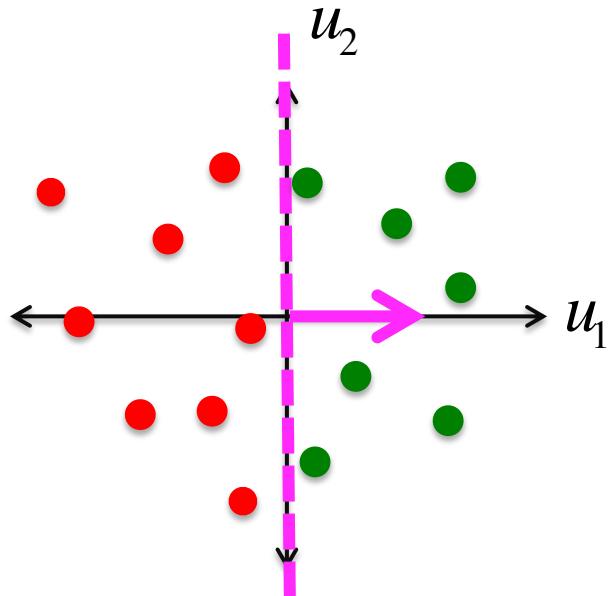
- This is the set of all vectors u that have zero projection along w .

All vectors on a line going through the origin and perpendicular to w !



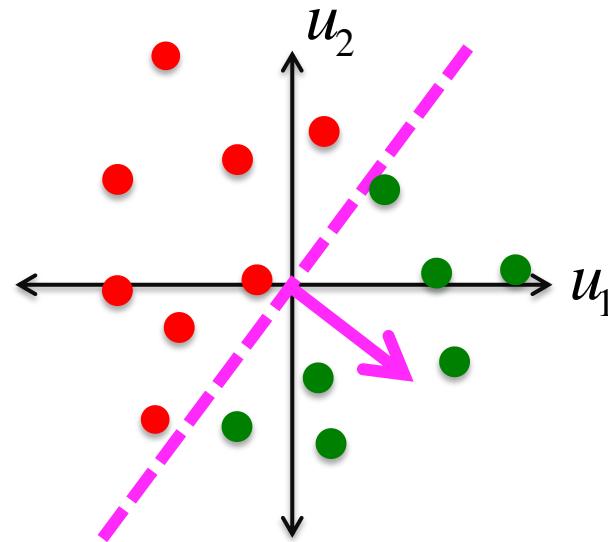
Classification in two dimensions

- Let's look at this for a few simple cases in two dimensions



$$\vec{w} = (1, 0)$$

$$\theta = 0$$



$$\vec{w} = (1, -1)$$

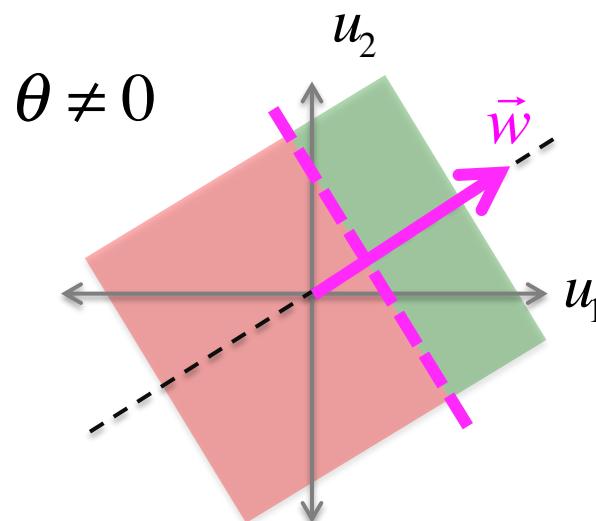
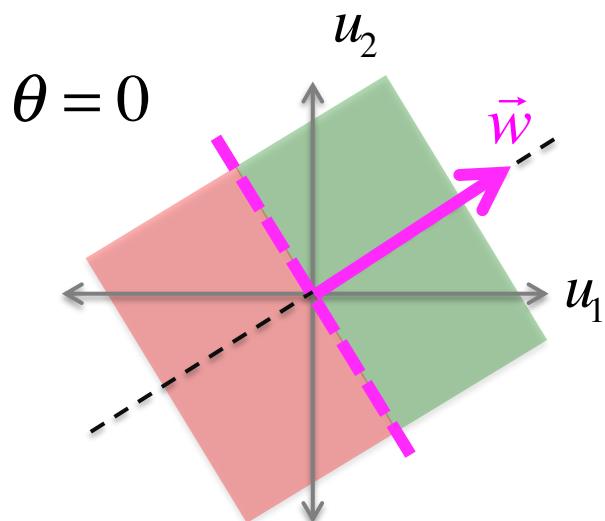
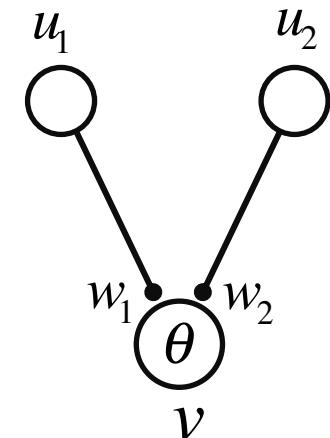
$$\theta = 0$$

Classification in two dimensions

- Now let's look at the case where $\theta \neq 0$

$$v = F(\vec{w} \cdot \vec{u} - \theta)$$

- Now the decision boundary is $\vec{w} \cdot \vec{u} = \theta$
- This is the set of all vectors \vec{u} whose projection along \vec{w} is given by θ .

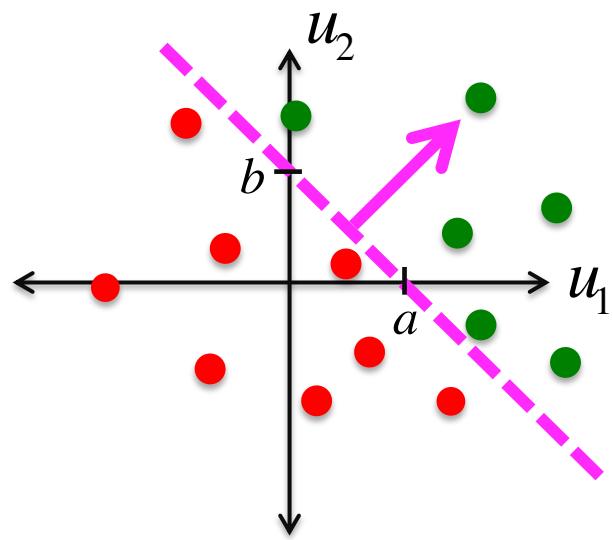


Classification in two dimensions

$$v = F(\vec{w} \cdot \vec{u} - \theta)$$

The decision boundary is $\vec{w} \cdot \vec{u} = \theta$

- Let's calculate the weight vector $\vec{w} = (w_1, w_2)$ that gives us the decision boundary shown below. Assume $\theta = 1$.



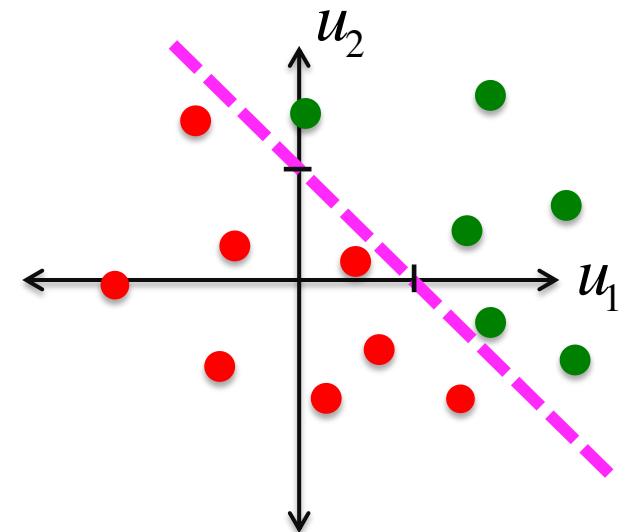
We have two points on the decision boundary we know, and two unknowns...

$$\vec{u}_a = (a, 0) \quad \vec{u}_a \cdot \vec{w} = \theta$$

$$\vec{u}_b = (0, b) \quad \vec{u}_b \cdot \vec{w} = \theta$$

Learning classification in higher dimensions

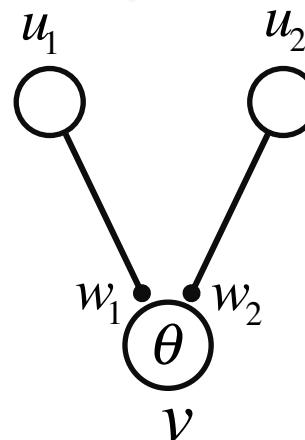
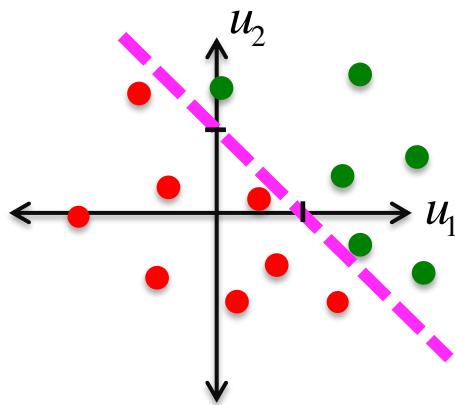
- In two dimensions, you can basically look at the data and decide where the decision boundary should be.
- But in higher dimensions this is a hard problem.



Perceptron learning

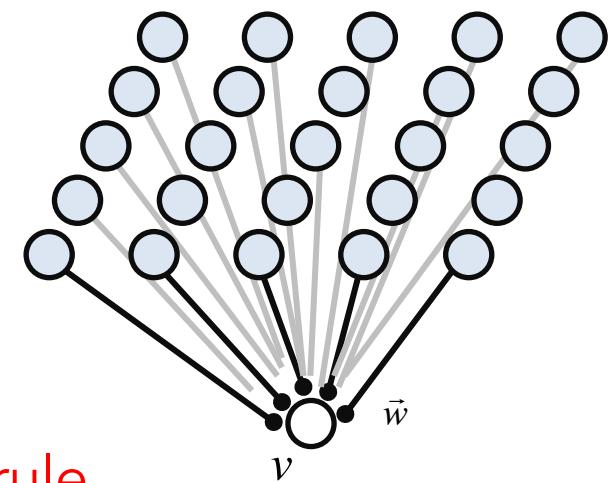
- How would we find the weight vector w that separates images of dogs from images of cats?

Low-dimensional



Perceptron learning rule

High-dimensional



Learning Objectives for Lecture 14

- Derive a mathematically tractable model of neural networks
(the rate model)
- Building receptive fields with neural networks
- Vector notation and vector algebra
- Neural networks for classification
- Perceptrons