


Introduction to Neural Computation

Prof. Michale Fee
MIT BCS 9.40 — 2017

Lecture 16
Networks, Matrices and Basis Sets

Seeing in high dimensions



<https://research.googleblog.com/2016/12/open-sourcing-embedding-projector-tool.html>

Learning Objectives for Lecture 16

- More on two-layer feed-forward networks
- Matrix transformations (rotated transformations)
- Basis sets
- Linear independence
- Change of basis

Learning Objectives for Lecture 16

- More on two-layer feed-forward networks
- Matrix transformations (rotated transformations)
- Basis sets
- Linear independence
- Change of basis

Two-layer feed-forward network

- We can expand our set of output neurons to make a more general network...


input firing rates

$$\left[u_1, u_2, u_3, \dots, u_{n_b} \right] = \vec{u}$$

Lots of synaptic weights! W_{ab}


output firing rates

$$\left[v_1, v_2, v_3, \dots, v_{n_a} \right] = \vec{v}$$



Two-layer feed-forward network

- We now have a weight from each of our input neurons onto each of our output neurons!
- We write the weights as a matrix.



weight matrix

$$W_{ab} = \begin{matrix} & \begin{matrix} b = 1 & 2 & 3 \end{matrix} \\ \begin{matrix} a = 1 \\ 2 \\ 3 \end{matrix} & \left[\begin{matrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \end{matrix} \right] \end{matrix}$$


a b
row column
post pre

$$= \left[\begin{matrix} \vec{w}_{a=1} \\ \vec{w}_{a=2} \\ \vec{w}_{a=3} \end{matrix} \right]$$

Two-layer feed-forward network

- We can write down the firing rates of our output neurons as a matrix multiplication.

$$\vec{v} = W \vec{u} \quad v_a = \sum_b W_{ab} u_b$$



$$\vec{v} = \begin{bmatrix} v_1 \\ v_2 \\ v_3 \end{bmatrix} = \begin{bmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \end{bmatrix} \begin{bmatrix} u_1 \\ u_2 \\ u_3 \end{bmatrix} = \begin{bmatrix} \vec{w}_{a=1} \cdot \vec{u} \\ \vec{w}_{a=2} \cdot \vec{u} \\ \vec{w}_{a=3} \cdot \vec{u} \end{bmatrix}$$

- Dot product interpretation of matrix multiplication

Two-layer feed-forward network

- There is another way to think about what the weight matrix means...

$$\vec{V} = W \vec{u} = \begin{bmatrix} b = 1 & 2 & 3 \\ w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \end{bmatrix} \begin{bmatrix} u_1 \\ u_2 \\ u_3 \end{bmatrix}$$

$\vec{w}^{(1)} \mid \vec{w}^{(2)} \mid \vec{w}^{(3)}$


vector of weights from input neuron 1 vector of weights from input neuron 2 vector of weights from input neuron 3

$$W = \begin{bmatrix} 1 & 0 & 1 \\ 1 & 1 & 1 \\ 0 & 0 & 1 \end{bmatrix}$$

Two-layer feed-forward network

- There is another way to think about what the weight matrix means...

$$\vec{v} = W \vec{u} = \begin{bmatrix} b = 1 & 2 & 3 \\ w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \\ \vec{w}^{(1)} & \vec{w}^{(2)} & \vec{w}^{(3)} \end{bmatrix} \begin{bmatrix} u_1 \\ u_2 \\ u_3 \end{bmatrix}$$




- What is the output if only input neuron 1 is active?

$$\vec{v} = \begin{bmatrix} w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \end{bmatrix} \begin{bmatrix} u_1 \\ 0 \\ 0 \end{bmatrix} = u_1 \begin{bmatrix} w_{11} \\ w_{21} \\ w_{31} \end{bmatrix} = u_1 \vec{w}^{(1)} = u_1 \begin{bmatrix} 1 \\ 1 \\ 0 \end{bmatrix}$$

Two-layer feed-forward network

$$\vec{v} = W \vec{u} = \begin{bmatrix} b = 1 & 2 & 3 \\ w_{11} & w_{12} & w_{13} \\ w_{21} & w_{22} & w_{23} \\ w_{31} & w_{32} & w_{33} \end{bmatrix} \begin{bmatrix} u_1 \\ u_2 \\ u_3 \end{bmatrix}$$

$$= \begin{bmatrix} \vec{w}^{(1)} & | & \vec{w}^{(2)} & | & \vec{w}^{(3)} \end{bmatrix}$$



$$\vec{v} = u_1 \begin{bmatrix} w_{11} \\ w_{21} \\ w_{31} \end{bmatrix} + u_2 \begin{bmatrix} w_{12} \\ w_{22} \\ w_{32} \end{bmatrix} + u_3 \begin{bmatrix} w_{13} \\ w_{23} \\ w_{33} \end{bmatrix}$$

$$\vec{v} = u_1 \vec{w}^{(1)} + u_2 \vec{w}^{(2)} + u_3 \vec{w}^{(3)}$$


$$W = \begin{bmatrix} 1 & 0 & 1 \\ 1 & 1 & 1 \\ 0 & 0 & 1 \end{bmatrix}$$

The output pattern is a linear combination of contributions from each of the input neurons!

Examples of simple networks

- Each input neuron connects to one neuron in the output layer, with a weight of one.

$$W = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix} \quad W = I$$



$$\vec{v} = W \vec{u} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} u_1 \\ u_2 \\ u_3 \end{bmatrix} = \begin{bmatrix} u_1 \\ u_2 \\ u_3 \end{bmatrix}$$

$$\vec{v} = \vec{u}$$

Examples of simple networks

- Each input neuron connects to one neuron in the output layer, with an arbitrary weight


$$W = \Lambda \quad \quad \Lambda = \begin{pmatrix} \lambda_1 & 0 & 0 \\ 0 & \lambda_2 & 0 \\ 0 & 0 & \lambda_3 \end{pmatrix}$$

$$\vec{v} = \Lambda \vec{u} = \begin{bmatrix} \lambda_1 & 0 & 0 \\ 0 & \lambda_2 & 0 \\ 0 & 0 & \lambda_3 \end{bmatrix} \begin{bmatrix} u_1 \\ u_2 \\ u_3 \end{bmatrix} = \begin{bmatrix} \lambda_1 u_1 \\ \lambda_2 u_2 \\ \lambda_3 u_3 \end{bmatrix}$$

Examples of simple networks

- Input neurons connect to output neurons with a weight matrix that corresponds to a rotation matrix.

$$W = \Phi \quad \Phi = \begin{bmatrix} \cos\phi & -\sin\phi \\ \sin\phi & \cos\phi \end{bmatrix}$$




$$\vec{v} = \Phi \cdot \vec{u} = \begin{bmatrix} \cos\phi & -\sin\phi \\ \sin\phi & \cos\phi \end{bmatrix} \begin{bmatrix} u_1 \\ u_2 \end{bmatrix} = \begin{bmatrix} u_1 \cos\phi - u_2 \sin\phi \\ u_1 \sin\phi + u_2 \cos\phi \end{bmatrix}$$

Examples of simple networks


- Let's look at an example rotation matrix ($\phi=-45^\circ$)

$$\Phi(-45^\circ) = \begin{bmatrix} \cos(-\frac{\pi}{4}) & -\sin(-\frac{\pi}{4}) \\ \sin(-\frac{\pi}{4}) & \cos(-\frac{\pi}{4}) \end{bmatrix} = \begin{bmatrix} \frac{1}{\sqrt{2}} & \frac{1}{\sqrt{2}} \\ -\frac{1}{\sqrt{2}} & \frac{1}{\sqrt{2}} \end{bmatrix}$$



$$\vec{v} = \Phi \cdot \vec{u} = \begin{bmatrix} \frac{1}{\sqrt{2}} & \frac{1}{\sqrt{2}} \\ -\frac{1}{\sqrt{2}} & \frac{1}{\sqrt{2}} \end{bmatrix} \begin{bmatrix} u_1 \\ u_2 \end{bmatrix}$$


$$\vec{v} = \frac{1}{\sqrt{2}} \begin{bmatrix} u_2 + u_1 \\ u_2 - u_1 \end{bmatrix}$$





Examples of simple networks

- Rotation matrices can be very useful when different directions in feature space carry different useful information

$$\Phi(-45^\circ) = \frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 1 \\ -1 & 1 \end{bmatrix}$$




Φ



Examples of simple networks

- Rotation matrices can be very useful when different directions in feature space carry different useful information




Learning Objectives for Lecture 16

- More on two-layer feed-forward networks
- Matrix transformations (rotated transformations)
- Basis sets
- Linear independence
- Change of basis

Matrix transformations

- In general A maps the set of vectors in \mathbb{R}^2 onto another set of vectors in \mathbb{R}^2 .


$$\vec{y} = A\vec{x}$$



Matrix transformations

- In general A maps the set of vectors in \mathbb{R}^2 onto another set of vectors in \mathbb{R}^2 .

$$\vec{y} = A\vec{x}$$




$$\vec{x} = A^{-1}\vec{y}$$

Matrix transformations


$$\vec{y} = A\vec{x}$$

- Perturbations from the identity matrix


$$A = \begin{pmatrix} 1+\delta & 0 \\ 0 & 1+\delta \end{pmatrix}$$




$$A = \begin{pmatrix} 1+\delta & 0 \\ 0 & 1 \end{pmatrix}$$




$$A = \begin{pmatrix} 1 & 0 \\ 0 & 1+\delta \end{pmatrix}$$




$$A = \begin{pmatrix} 1+\delta & 0 \\ 0 & 1-\delta \end{pmatrix}$$



$$A = \begin{pmatrix} -1 & 0 \\ 0 & 1 \end{pmatrix}$$



$$A = \begin{pmatrix} -1 & 0 \\ 0 & -1 \end{pmatrix}$$



These are all diagonal matrices

$$\Lambda = \begin{pmatrix} a & 0 \\ 0 & b \end{pmatrix}$$

$$\Lambda^{-1} = \begin{pmatrix} a^{-1} & 0 \\ 0 & b^{-1} \end{pmatrix}$$

Rotation matrix

- Rotation in 2 dimensions





$$\Phi(\theta) = \begin{pmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{pmatrix}$$

$$\theta = 10^\circ$$

$$\theta = 25^\circ$$

$$\theta = 45^\circ$$

$$\theta = 90^\circ$$




- Does a rotation matrix have an inverse? $\det(\Phi) = 1$
- The inverse of a rotation matrix is just its transpose

$$\Phi^{-1}(\theta) = \Phi(-\theta) = \Phi^T(\theta)$$

Rotated transformations

- Let's construct a matrix that produces a stretch along a 45° angle...

$$\Phi = \begin{pmatrix} \cos 45 & -\sin 45 \\ \sin 45 & \cos 45 \end{pmatrix} = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & -1 \\ 1 & 1 \end{pmatrix}$$




- We do each of these steps by multiplying our matrices together

$$\Phi \Lambda \Phi^T \vec{x}$$

Rotated transformations

- Let's construct a matrix that produces a stretch along a 45° angle...



$$\begin{aligned}\Phi^T &= \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ -1 & 1 \end{pmatrix} & \Lambda &= \begin{pmatrix} 2 & 0 \\ 0 & 1 \end{pmatrix} & \Phi &= \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & -1 \\ 1 & 1 \end{pmatrix}\end{aligned}$$

$$\Phi\Lambda\Phi^T = \frac{1}{2} \begin{pmatrix} 3 & 1 \\ 1 & 3 \end{pmatrix}$$

Inverse of matrix products

- We can undo our transformation by taking the inverse

$$[\Phi \Lambda \Phi^T]^{-1}$$

- How do you take the inverse of a sequence of matrix multiplications A^*B^*C ?

$$[ABC]^{-1} = C^{-1}B^{-1}A^{-1}$$

$$\begin{aligned}[ABC]^{-1}ABC &= C^{-1}B^{-1}A^{-1}ABC \\ &= C^{-1}B^{-1}BC\end{aligned}$$

- Thus...

$$[\Phi \Lambda \Phi^T]^{-1} = [\Phi^T]^{-1} [\Lambda]^{-1} [\Phi]^{-1}$$

$$= C^{-1}C$$


$$= I$$

$$[\Phi \Lambda \Phi^T]^{-1} = \Phi \Lambda^{-1} \Phi^T$$

$$\Lambda^{-1} = \begin{pmatrix} 2 & 0 \\ 0 & 1 \end{pmatrix}^{-1} = \begin{pmatrix} 1/2 & 0 \\ 0 & 1 \end{pmatrix}$$

Rotated transformations

- Let's construct a matrix that undoes a stretch along a 45° angle...



$$\Phi^T$$

$$= \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ -1 & 1 \end{pmatrix}$$

$$\Lambda^{-1}$$

$$= \begin{pmatrix} 1/2 & 0 \\ 0 & 1 \end{pmatrix}$$


$$\Phi(+45^\circ)$$

$$= \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & -1 \\ 1 & 1 \end{pmatrix}$$

$$\Phi\Lambda^{-1}\Phi^T = \frac{1}{4} \begin{pmatrix} 3 & -1 \\ -1 & 3 \end{pmatrix}$$

Rotated transformations

- Construct a matrix that does compression along a -45° angle...



Φ^T

Λ

$\Phi(-45^\circ)$

$$= \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & -1 \\ 1 & 1 \end{pmatrix}$$

$$= \begin{pmatrix} 0.2 & 0 \\ 0 & 1 \end{pmatrix}$$

$$= \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ -1 & 1 \end{pmatrix}$$

$$\Phi \Lambda \Phi^T = \begin{pmatrix} 0.6 & 0.4 \\ 0.4 & 0.4 \end{pmatrix}$$

Transformations that can't be undone




- Some transformation matrices have no inverse...

$$\vec{x}$$

$$\Phi^T \vec{x}$$

$$\Lambda \Phi^T \vec{x}$$

$$\Phi \Lambda \Phi^T \vec{x}$$



$$\Phi^T = \Phi(45^\circ)$$

$$\Lambda$$

$$\Phi(-45^\circ)$$

$$= \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & -1 \\ 1 & 1 \end{pmatrix}$$

$$= \begin{pmatrix} 0 & 0 \\ 0 & 1 \end{pmatrix}$$

$$= \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ -1 & 1 \end{pmatrix}$$

$$\Phi \Lambda \Phi^T = \begin{pmatrix} 0.5 & 0.5 \\ 0.5 & 0.5 \end{pmatrix}$$

$$\det(\Phi \Lambda \Phi^T) = 0$$

$$\det(\Lambda) = 0$$

Learning Objectives for Lecture 16

- More on two-layer feed-forward networks
- Matrix transformations (rotated transformations)
- **Basis sets**
- Linear independence
- Change of basis

Basics of basis sets

- We can think of vectors as abstract 'directions' in space. But in order to specify the elements of a vector, we need to choose a coordinate system.
- To do this, we write our vector as a linear combination of a set of special vectors called the 'basis set.'

$$\vec{v} = \begin{pmatrix} x \\ y \\ z \end{pmatrix} = x \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix} + y \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix} + z \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix} = x\hat{e}_1 + y\hat{e}_2 + z\hat{e}_3$$

- The numbers x, y, z are called the coordinates of the vector.
- The vectors $\{\hat{e}_1, \hat{e}_2, \hat{e}_3\}$ are called the 'basis vectors', in this case, in three dimensions.

Basics of basis sets

- In order to describe an arbitrary vector in the space of real numbers in n dimensions (\mathbb{R}^n), our basis vectors need to have n numbers.
- In order to describe an arbitrary vector in \mathbb{R}^n , we need to have n basis vectors.
- The basis set we wrote down earlier $\{\hat{e}_1, \hat{e}_2, \hat{e}_3\}$ is called the 'standard basis'. Each vector has one element that's a one and the rest are zeros.

$$\hat{e}_1 = \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix} \quad \hat{e}_2 = \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix} \quad \hat{e}_3 = \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix}$$

Orthonormal basis

- In addition, the standard basis has the interesting property that

each vector is a unit vector

$$\hat{e}_i \cdot \hat{e}_i = 1$$

$$\hat{e}_1 = \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix} \quad \hat{e}_2 = \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix} \quad \hat{e}_3 = \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix}$$

- Each vector is orthogonal to all the other vectors

$$\hat{e}_1 \cdot \hat{e}_2 = 0 \quad \hat{e}_1 \cdot \hat{e}_3 = 0 \quad \hat{e}_2 \cdot \hat{e}_3 = 0 \quad \hat{e}_i \cdot \hat{e}_j = 0, \quad i \neq j$$

- These properties can be written compactly as

$$\hat{e}_i \cdot \hat{e}_j = \delta_{ij} \quad \delta_{ij} = \begin{cases} 1 & \text{if } i = j \\ 0 & \text{if } i \neq j \end{cases}$$

- Any basis set with these properties is called 'orthonormal'.


Basics of basis sets

- The standard basis is not the only orthonormal basis

Consider a different set of orthogonal unit vectors: $\{\vec{f}_1, \vec{f}_2\}$

$$\vec{v} = (\vec{v} \cdot \hat{f}_1) \hat{f}_1 + (\vec{v} \cdot \hat{f}_2) \hat{f}_2$$


$$\vec{v}_f = \begin{pmatrix} \vec{v} \cdot \hat{f}_1 \\ \vec{v} \cdot \hat{f}_2 \end{pmatrix}$$



- The vector coordinates are given by the dot products of the vector \vec{v} with each of the basis vectors.

Non-orthonormal basis sets

- Vectors can also be written as a linear combination of (almost) any vectors, not just orthonormal basis vectors



Basics of basis sets

- Let's decompose an arbitrary vector v in a basis set $\{\vec{f}_1, \vec{f}_2\}$

$$\vec{v} = c_1 \vec{f}_1 + c_2 \vec{f}_2$$

- The coefficients c_1 and c_2 are called 'coordinates of the vector v in the basis $\{\vec{f}_1, \vec{f}_2\}$ '.

- The vector $\vec{v}_f = \begin{pmatrix} c_1 \\ c_2 \end{pmatrix}$ is called the 'coordinate vector' of \vec{v} in the basis $\{\vec{f}_1, \vec{f}_2\}$.

Basics of basis sets

- Let's look at an example. Consider the basis

$$\{\vec{f}_1, \vec{f}_2\} = \left\{ \begin{pmatrix} 1 \\ 3 \end{pmatrix}, \begin{pmatrix} -2 \\ 1 \end{pmatrix} \right\}$$

and the vector $\vec{v} = \begin{pmatrix} 3 \\ 5 \end{pmatrix}$ in the standard basis.

- Find the vector coordinates of \vec{v} in the new basis.
- Write \vec{v} as a linear combination of the new basis vectors:

$$c_1 \vec{f}_1 + c_2 \vec{f}_2 = \vec{v}$$

system of equations

$$c_1 \begin{pmatrix} 1 \\ 3 \end{pmatrix} + c_2 \begin{pmatrix} -2 \\ 1 \end{pmatrix} = \begin{pmatrix} 3 \\ 5 \end{pmatrix}$$
$$\begin{aligned} c_1 - 2c_2 &= 3 \\ 3c_1 + c_2 &= 5 \end{aligned}$$

Basics of basis sets

- We can write this system of equations in matrix notation:

$$c_1 - 2c_2 = 3$$

$$3c_1 + c_2 = 5$$

$$F \vec{v}_f = \vec{v}$$

where $F = \begin{pmatrix} 1 & -2 \\ 3 & 1 \end{pmatrix}$ $\vec{v}_f = \begin{pmatrix} c_1 \\ c_2 \end{pmatrix}$ $\vec{v} = \begin{pmatrix} 3 \\ 5 \end{pmatrix}$

- Now solve for \vec{v}_f by multiplying both sides of the equation by the inverse of matrix F .

$$F^{-1} F \vec{v}_f = F^{-1} \vec{v}$$

$$\vec{v}_f = F^{-1} \vec{v}$$

Basics of basis sets

- We can find the inverse of F :

$$F^{-1} = \frac{1}{7} \begin{pmatrix} 1 & 2 \\ -3 & 1 \end{pmatrix}$$

$$\vec{v}_f = F^{-1} \vec{v} = \frac{1}{7} \begin{pmatrix} 1 & 2 \\ -3 & 1 \end{pmatrix} \begin{pmatrix} 3 \\ 5 \end{pmatrix}$$

$$= \frac{1}{7} \begin{pmatrix} 3+10 \\ -9+5 \end{pmatrix} = \frac{1}{7} \begin{pmatrix} 13 \\ -4 \end{pmatrix}$$

- Thus, we find the coordinate vector of v in basis $\{\vec{f}_1, \vec{f}_2\}$

$$\vec{v}_f = \begin{pmatrix} 13/7 \\ -4/7 \end{pmatrix}$$

Basics of basis sets

- In summary: to find the coordinate vector for v in the basis $\{\vec{f}_1, \vec{f}_2\}$, we construct a matrix F whose columns are just the elements of the basis vectors.

$$F = \left(\begin{array}{c|c} \vec{f}_1 & \vec{f}_2 \end{array} \right)$$

$$F = \left(\begin{array}{c|c|c|c} \vec{f}_1 & \vec{f}_2 & \vec{f}_3 & \dots & \vec{f}_n \end{array} \right)$$

such that $\vec{v} = F \vec{v}_f$

- We can solve for \vec{v}_f by multiplying both sides of the equation by the inverse of matrix F

$$\vec{v}_f = F^{-1} \vec{v} \quad \text{'change of basis'}$$

- But this only works if F has an inverse!

Learning Objectives for Lecture 16

- More on two-layer feed-forward networks
- Matrix transformations (rotated transformations)
- Basis sets
- Linear independence
- Change of basis

Subspaces

- We need n vectors in \mathbb{R}^n to form a basis in \mathbb{R}^n . But not any set of n vectors will do the trick!
- Consider the following set of vectors

$$\vec{f}_1 = \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix} \quad \vec{f}_2 = \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix} \quad \vec{f}_3 = \begin{pmatrix} 1 \\ 1 \\ 0 \end{pmatrix}$$

- Note that any linear combination of $\{\vec{f}_1, \vec{f}_2, \vec{f}_3\}$ will always lie in the (x, y) plane

$$\vec{v} = c_1 \vec{f}_1 + c_2 \vec{f}_2 + c_3 \vec{f}_3 = \begin{pmatrix} c_1 + c_3 \\ c_2 + c_3 \\ 0 \end{pmatrix}$$

- Thus, the set of vectors $\{\vec{f}_1, \vec{f}_2, \vec{f}_3\}$ doesn't span all of \mathbb{R}^3
It only spans the x - y plane - a subspace of \mathbb{R}^3

Linear independence

$$\vec{f}_1 = \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix} \quad \vec{f}_2 = \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix} \quad \vec{f}_3 = \begin{pmatrix} 1 \\ 1 \\ 0 \end{pmatrix}$$

- Note that we can write any of these vectors as a linear combination of the other two.

$$\vec{f}_3 = \vec{f}_1 + \vec{f}_2$$

$$\vec{f}_2 = \vec{f}_3 - \vec{f}_1$$


$$\vec{f}_1 = \vec{f}_3 - \vec{f}_2$$

- Thus, this set of vectors is called 'linearly dependent'.
- Any set of n linearly dependent vectors cannot form a basis in \mathbb{R}^n
- How do you know if a set of vectors is linearly dependent?

$$F = \left(\begin{array}{c|c|c|c} \vec{f}_1 & \vec{f}_2 & \vec{f}_3 & \dots & \vec{f}_n \end{array} \right) \quad \det(F) = 0$$

Linear dependence

- If $\det(F) = 0$ then F maps \vec{v}_f into a subspace of \mathbb{R}^n




- If F maps onto a subspace, then the mapping is not reversible!


$$\det(F) = 0$$

Geometric interpretation of determinant

- The determinant is the ‘volume’ of a unit cube after transformation (area of unit square in two dimensions).



- A pure rotation matrix has a determinant of one.



Learning Objectives for Lecture 16

- More on two-layer feed-forward networks
- Matrix transformations (rotated transformations)
- Basis sets
- Linear independence
- Change of basis

Change of basis

$$\left\{ \vec{f}_1, \vec{f}_2, \dots, \vec{f}_n \right\} \quad F = \left(\begin{array}{c|c|c|c} \vec{f}_1 & \vec{f}_2 & \dots & \vec{f}_n \end{array} \right)$$

- If $\det(F) \neq 0$ then the vectors $\left\{ \vec{f}_1, \vec{f}_2, \dots, \vec{f}_n \right\}$
 - are linearly independent
 - form a complete basis set in \mathbb{R}^n
- Then the matrix F implements a ‘change of basis’

From standard basis to \vec{f}

$$\vec{v}_f = F^{-1} \vec{v}$$

Or from \vec{f} to standard basis

$$\vec{v} = F \vec{v}_f$$

Change of basis

- The change of basis is easy if $\{\vec{f}_1, \vec{f}_2\}$ is an orthonormal basis...

$$F = \begin{pmatrix} | & | \\ \hat{f}_1 & \hat{f}_2 \\ | & | \end{pmatrix} \quad F^T = \begin{pmatrix} - & \hat{f}_1 & - \\ - & \hat{f}_2 & - \end{pmatrix}$$

$$F^T F = \begin{pmatrix} - & \hat{f}_1 & - \\ - & \hat{f}_2 & - \end{pmatrix} \begin{pmatrix} | & | \\ \hat{f}_1 & \hat{f}_2 \\ | & | \end{pmatrix} = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix} = I$$

Thus...

$$F^T = F^{-1}$$

F is just a rotation matrix!

Change of basis

- With an orthonormal basis set, the coordinates are just given by the dot product with the basis vectors !

$$F = \begin{pmatrix} | & | \\ \hat{f}_1 & \hat{f}_2 \\ | & | \end{pmatrix} \quad F^{-1} = F^T = \begin{pmatrix} -\hat{f}_1- \\ -\hat{f}_2- \end{pmatrix}$$

$$\vec{v}_f = F^{-1}\vec{v}$$

$$\vec{v}_f = F^T \vec{v} = \begin{pmatrix} -\hat{f}_1- \\ -\hat{f}_2- \end{pmatrix} \vec{v} = \begin{pmatrix} \vec{v} \cdot \hat{f}_1 \\ \vec{v} \cdot \hat{f}_2 \end{pmatrix}$$


Change of basis

- In two dimensions, there is a family of orthonormal basis sets

$$\hat{f}_1 = \begin{pmatrix} \cos\theta \\ -\sin\theta \end{pmatrix} \quad \hat{f}_2 = \begin{pmatrix} \sin\theta \\ \cos\theta \end{pmatrix} \quad F = \begin{pmatrix} \cos\theta & \sin\theta \\ -\sin\theta & \cos\theta \end{pmatrix}$$



$$\vec{v} = (\vec{v} \cdot \hat{f}_1) \hat{f}_1 + (\vec{v} \cdot \hat{f}_2) \hat{f}_2$$

$$\vec{v}_f = F^T \vec{v} \quad \vec{v}_f = \begin{pmatrix} \vec{v} \cdot \hat{f}_1 \\ \vec{v} \cdot \hat{f}_2 \end{pmatrix}$$



- The vector coordinates are given by the dot products of the vector \vec{v} with each of the rotated basis vectors.

Seeing in high dimensions



<https://research.googleblog.com/2016/12/open-sourcing-embedding-projector-tool.html>

Learning Objectives for Lecture 16

- More on two-layer feed-forward networks
- Matrix transformations (rotated transformations)
- Basis sets
- Linear independence
- Change of basis